

References

- [Training LLM's to Reason in Continuous a Latent Space](https://arxiv.org/pdf/2412.06769)
(<https://arxiv.org/pdf/2412.06769>)
- [From Explicit CoT to Implicit CoT: Learning to Internalize CoT Step by Step](https://arxiv.org/pdf/2405.14838)
(<https://arxiv.org/pdf/2405.14838>)
 - simply deletes CoT steps rather than replacing them with continuous thoughts
 - the deleted steps become "internalized" in the model's parameters
 - gradually delete steps until all are deleted and all CoT is internalized
 - [github](https://github.com/da03/Internalize_CoT_Step_by_Step) ([https://github.com/da03/Internalize CoT Step by Step](https://github.com/da03/Internalize_CoT_Step_by_Step))

Reasoning in Latent Space

Recall the behavior of LLM's that use reasoning/"Chain of Thought"

Given prompt

$\backslash \mathbf{x}_{(1:\bar{T})}$

rather than *immediately* producing response

$\backslash \mathbf{y}_{(1:\bar{T})}$

giving computation trace

$\backslash \mathbf{x}, \backslash \mathbf{y}$

- (dropping the sequence subscript that indexes tokens)

an LLM is trained to think "Step by Step"

- creating a sequence of steps
- the Chain of Thought ("reasoning" trace) $\backslash \mathbf{rat}$
- enumerating sequential steps of a process that produces the response

$$\backslash \mathbf{rat} = [\backslash \mathbf{rat}_{(1)}, \dots, \backslash \mathbf{rat}_{(\text{num_thoughts})}]$$

- where each $\backslash \mathbf{rat}_{\text{tp}}$ is a thought represented as multi-token sequence

All the tokens in sequences $\backslash \mathbf{x}$, $\backslash \mathbf{rat}$, $\backslash \mathbf{y}$ come from the vocabulary $\backslash \mathbf{Vocab}$.

Thus, all the token sequences are in "Natural Language".


But the reasoning trace $\backslash \mathbf{rat}$ is often hidden from the user.

- so $\backslash \mathbf{rat}$ is not visible to the user

Does the LLM have to "think" in Natural Language ?

This paper proposes adapting CoT

Chain of Thought



Attribution: <https://arxiv.org/pdf/2412.06769#page=2>

(Note: paper uses $\backslash \mathbf{x}$ as a denotation of the output sequence)

to include reasoning tokens

- that are not restricted to discrete categorical values
 - "Natural Language" tokens representing elements of vocabulary
 \ **Vocab**
- but are continuous vectors

Being continuous rather than discrete

- enables a richer space of "thoughts"

This is called *Chain of Continuous Thought (Coconut)*.

Chain of Thought




Figure 1 A comparison of Chain of Continuous Thought (CoCONUT) with Chain-of-Thought (CoT). In CoT, the model generates the reasoning process as a word token sequence (e.g., $[x_i, x_{i+1}, \dots, x_{i+j}]$ in the figure). CoCONUT regards the last hidden state as a representation of the reasoning state (termed “continuous thought”), and directly uses it as the next input embedding. This allows the LLM to reason in an unrestricted latent space instead of a language space.

Attribution: <https://arxiv.org/pdf/2412.06769#page=2>

Continuous thoughts


Where do the tokens that form the continuous thoughts come from ?

It is a simple adaptation of the Transformer architecture.


- see diagram below

In [5]: fig_transf

Out[5]:



Latent state



The Transformer operates in Auto-regressive mode:

- outputs one token at a time
- at step t its output $\hat{y}_t$
- is appended to the previous output sequence $\hat{y}_{(1:t-1)}$
- which becomes the input for step $t+1$

The tokens (vectors of length $|\mathcal{V}|$, where $\mathcal{V}$ is the Vocabulary of tokens) are

- *embedded* into vectors of length $d = d_{\text{model}}$
- and processed by the "body" of the Transformer
 - everything before the Linear layer

The output before the linear layer (referred to as the "Latent state")

- is a vector of length d
- which is *un-embedded* by the Linear layer into a vector of length $|\mathcal{V}|$
 - a probability vector over the finite, discrete set of language tokens $\mathcal{V}$
 - so a single token in $\mathcal{V}$ can be output

We represent this schematic description of the Transformer in the left plot "Language Mode".

The Embed block

- transforms discrete tokens into vectors whose length d is the path width of the Transformer


The Un-embed block

- transforms the latent state vector of length d back into a Language token
 - which must then be embedded before it can be used in the next time step


In [7]: fig_combined

Out[7]:

Transformer Schematic: Language Mode



Transformer Schematic: Latent Mode



The Continuous Thought tokens

- are the final vectors of the Body
 - length d
 - *Latent state*

They are *not* un-embedded.

They can be fed *directly* into the Body at the next time step

- by-passing the Un-embed of step and the subsequent Embed of step (+1).

See the schematic diagram for "Latent Mode"

What does a Continuous token represent ?

Since a continuous token is not output to the user, it's vector representation does not have to be interpretable.

But the authors conducted an experiment

- decode the continuous token using the Un-embed

Here is the input to the LLM

```
James decides to run 3 sprints 3 times a week.  
He runs 60 meters each sprint.  
How many total meters does he run a week ?
```

There are several paths that lead to the correct answer ($3 * 3 * 60 = 540$)

- First compute ($3 * 3$), then multiple the result by 60
 - leading to outputting 9 as the first token of the reasoning trace
- First compute ($60 * 3$)
 - leading to outputting 180 as the first token of the reasoning trace

When the continuous thought vector is un-embedded

- it shows a probability distribution over the likely next tokens: 9 and 180

The authors interpret the thought vector as a "super-position" of possible next tokens.

Decoding a Continuous Thought




Figure 4 A case study where we decode the continuous thought into language tokens.

Attribution: <https://arxiv.org/pdf/2412.06769#page=7>

Absent "latent mode"

- the LLM would have to commit to **one** of the choices
 - e.g., "180" has the highest probability
 - in extending the reasoning trace
- by not committing immediately:
 - the LLM can "simultaneously" explore **both** possibilities
 - essentially allowing multiple traces (one beginning with "180", another beginning with "9")

The authors interpret this as the LLM

- performing a search across possible completions
- preserving possibilities
- until a later point in time

Training the LLM to reason in mixed language/latent mode

An LLM can be trained in CoT reasoning by using examples of reasoning traces

$\backslash \mathbf{x}$, $\backslash \mathbf{rat}$, $\backslash \mathbf{y}$

But this works only if the tokens of every thought in $\backslash \mathbf{rat}$ are *language* tokens.

How do we train when $\backslash \mathbf{rat}$ is a mixture of language and continuous tokens?

- the continuous tokens are *defined* by training; not pre-specified

Clarification

Trace \rat consists of

- a number of *steps*
- each step consists of c "thoughts" (tokens: either language or continuous)

The authors use a clever technique:

Multi-stage training *curriculum*

- different training set at each stage
- at stage k
 - replace the first k steps ($k * c$ tokens)
 - by continuous steps ($k * c$ continuous thoughts)
- stage 0 is traditional CoT training examples

At stage 0, the LLM learns how to reason.

- calculate a loss per language token in `\rat`
- guiding the learning-to-reason process

At subsequent stages, it learns to derive continuous tokens incrementally

- increasing k , the number of continuous thoughts

By proceeding incrementally

- the model does not "un-learn" traditional CoT in language tokens

The loss for continuous tokens

- is **masked** (not calculated)
 - no actual target with which to compare the prediction
- the loss of the language tokens following the continuous tokens
 - is reduced
 - when the preceding continuous tokens contribute to the prediction of the language tokens

The underlined elements in the diagram are the ones for which a loss is calculated.

Note

- each thought is a sequence of up to c tokens

Continuous tokens bracketed by `<bot>`, `<eot>`

Decoding a Continuous Thought




Figure 2 Training procedure of Chain of Continuous Thought (Coconut). Given training data with language reasoning steps, at each training stage we integrate c additional continuous thoughts ($c = 1$ in this example), and remove one language reasoning step. The cross-entropy loss is then used on the remaining tokens after continuous thoughts.

Attribution: <https://arxiv.org/pdf/2412.06769#page=4>

In [8]: `print("Done")`

Done

